

Contents

1	Games of Complete Information	2
1.1	Static Games of Complete Information	2
1.2	Dynamic Games of Complete Information	8
1.2.1	Perfect Information	8

Games of Complete Information

1.1 Static Games of Complete Information

We start with a classic story:

Prisoners' Dilemma

Two suspects are arrested and charged with a crime. The police lack sufficient evidence to convict the suspects, unless at least one confesses. The suspects are held in separate cells and told:

- If only one of you confesses and testifies against your partner, then the person who confesses will go free while the person who does not confess will be convicted and given a 20-year jail sentence.
- If both of you confess, then both will be convicted and sent to prison for 5 years.
- Finally, if neither of you confesses, then both of you will be convicted of a minor offence and sentenced to 1 year in jail.

What should the suspects do?

To investigate the situation, we can represent the prisoners' decisions using a matrix:

	Holdout	Confess
Holdout	$(-1, -1)$	$(-20, 0)$
Confess	$(0, -20)$	$(-5, -5)$

With some observation, it is easy to see that no matter what the other prisoner chooses to do, the best decision for the prisoner is to confess. Note that in this representation, we represent a *strategy* with an ordered pair consisting of the decisions made by the two prisoners respectively. Extending this idea to a higher dimension, we see that every strategy of a game is formed by an ordered n -tuple, where each factor represents a decision taken by the corresponding player.

Definition 1.1.1 ► Strategy Space

In an n -player game, the **strategy space** of the i -th player, denoted by S_i , is the set of all possible decisions for the player.

In the above example, for each strategy (i, j) , we then place the years of imprisonment, or the *payoff*, for the two prisoners following that strategy in the corresponding matrix entry. Intuitively, for an n -player game in general, we can assign a payoff value for each strategy in $\prod_{i=1}^n S_i$, the set of all strategies.

Definition 1.1.2 ► Payoff Function

Let S_i be the strategy space of the i -th player in an n -player game, the **payoff function** for the i -th player is a map $u_i : \prod_{j=1}^n S_j \rightarrow \mathbb{R}$.

This gives us a convenient way to represent a game quantitatively.

Definition 1.1.3 ► Normal-Form of Games

Let G be an n -player game. Let S_i be the strategy space and u_i be the payoff function of the i -th player, then the **normal-form**, or **strategic form** of G is defined as

$$G := \{S_1, S_2, \dots, S_n, u_1, u_2, \dots, u_n\}.$$

It is important to note that the payoff a player depends not only on his own action but also on the actions of others. This inter-dependence is the essence of games.

Remark. For 2-player games, the payoff can be represented in a matrix.

Given any game, our goal is to derive the optimal strategy for the players. Before that, we shall discuss some criteria under which a strategy cannot be optimal. For the sake of notational simplicity, we denote

$$\mathbf{s}_{-i} := (s_1, \dots, s_{i-1}, s_{i+1}, \dots, s_n)$$

and

$$\mathbf{s} := (s_1, s_2, \dots, s_n) = (s_i, \mathbf{s}_{-i}).$$

Analogously, we denote

$$S_{-i} := \prod_{j=1}^{i-1} S_j \times \prod_{j=i+1}^n S_j.$$

Definition 1.1.4 ► Dominated and Dominant Strategies

Let $G := \{S_1, S_2, \dots, S_n, u_1, u_2, \dots, u_n\}$ be a normal-form game and $s_i, s'_i \in S_i$ be two strategies for the i -th player. We say that s_i is **strictly dominated** by s'_i or that s'_i is

strictly dominant over s_i if

$$u_i(s_i, \mathbf{s}_{-i}) < u_i(s'_i, \mathbf{s}_{-i}) \quad \text{for all } \mathbf{s}_{-i} \in S_{-i}.$$

Remark. In other words, s_i is strictly dominated by s'_i if and only if it is a strictly worse decision no matter what the other players decide to do.

If a strategy s is strictly dominated by any other strategy, then it is not a rational strategy because it can never be optimal.

Assuming that in an n -player game, every player knows exactly what the strategies are and their respective payoffs, then we can proceed with the following reasoning:

Technique 1.1.5 ▶ Iterated Elimination of Strictly Dominated Strategies (IESDS)

- Suppose that s_{i_0} is strictly dominated by any other strategy of player i , then if the player is rational, the strategy will never be played.
- Therefore, we can reduce the strategy space to $\{\mathbf{s} \in \prod_{i=1}^n S_i : s_i \neq s_{i_0}\}$.
- Apply the same reasoning on the reduced strategy space iteratively. If at each iteration there exists such a strictly dominated strategy, then we can successfully obtain the optimal strategy for the n -players.

However, in most games, we might not be able to find such a strictly dominated strategy. Therefore, we have to think from a conditional context: given a set of strategies played by the other players, what is the *best response* for the current player?

Definition 1.1.6 ▶ Best-Response Correspondence

Let $G := \{S_1, S_2, \dots, S_n, u_1, u_2, \dots, u_n\}$ be a normal-form game. The **best-response correspondence** for player i is a map $R_i : S_{-i} \rightarrow \mathcal{P}(S_i)$ defined by

$$R_i(\mathbf{s}_{-i}) := \operatorname{argmax}_{s_i \in S_i} u_i(s_i, \mathbf{s}_{-i}).$$

Now, given $\mathbf{s}_{-i}$, the actions played by the other players, $R_i(\mathbf{s}_{-i})$ tells us what the best responses are for player i for his own profit. In other words, player i has **no incentive** to choose any strategy in $S_i \setminus R_i(\mathbf{s}_{-i})$. If such a situation holds for all the players, then it is clear that the game has reached an *equilibrium* where no player is incentivised to deviate from the current strategy.

It is clear that such equilibria occur exactly when the strategies of all players are the best responses of one another.

Definition 1.1.7 ▶ Nash Equilibrium

Let $G := \{S_1, S_2, \dots, S_n, u_1, u_2, \dots, u_n\}$ be a normal-form n -player game, then the strategies $(s_1^*, s_2^*, \dots, s_n^*)$ are a **Nash equilibrium** if

$$s_i^* \in R_i(s_{-i}^*) \quad \text{for all } i = 1, 2, \dots, n,$$

or equivalently,

$$u_i(s_i^*, s_{-i}^*) = \max_{s_i \in S_i} u_i(s_i, s_{-i}^*) \quad \text{for all } i = 1, 2, \dots, n.$$

Now, define

$$G(R_i) := \{(s_i, s_{-i}) : s_i \in R_i(s_{-i}), s_{-i} \in S_{-i}\}$$

as the **graph** of R_i , then one can check that $(s_1^*, s_2^*, \dots, s_n^*)$ are a Nash equilibrium if and only if

$$(s_1^*, s_2^*, \dots, s_n^*) \in \bigcap_{i=1}^n G(R_i).$$

Next, let us consider a simple case of bi-matrix games. Take any guess $(s_1, s_2) \in S_1 \times S_2$. We can compute $R_2(s_1)$ and $R_1(s_2)$, and (s_1, s_2) are a Nash equilibrium if and only if $s_1 \in R_1(s_2)$ and $s_2 \in R_2(s_1)$ by Definition 1.1.7.

Therefore, we can simply repeat this calculation for every pair $(s_1, s_2) \in S_1 \times S_2$ and deduce the set $G(R_1) \cap G(R_2)$, which consists of all Nash equilibria.

Proposition 1.1.8 ▶ Relation between Nash Equilibrium and IESDS

Let $G = (S_1, S_2, \dots, S_n, u_1, u_2, \dots, u_n)$ be an n -player normal-form game. If the strategies $(s_1^*, s_2^*, \dots, s_n^*)$ are a Nash equilibrium, then each s_i^* cannot be eliminated by IESDS.

Proof. Suppose on contrary that there exists some s_i^* which is eliminated by IESDS. Without loss of generality, assume s_i^* to be the first in $(s_1^*, s_2^*, \dots, s_n^*)$ to be eliminated. By Technique 1.1.5, there exists some $s'_i \in S_i$ with $s'_i \neq s_i^*$ such that

$$u_i(s_i^*, s_{-i}) < u_i(s'_i, s_{-i})$$

for all $s_{-i} \in S_{-i}$. In particular, this means that

$$u_i(s_i^*, s_{-i}^*) < u_i(s'_i, s_{-i}^*).$$

However, by Definition 1.1.7, we have $u_i(s_i^*, s_{-i}^*) = \max_{s_i \in S_i} u_i(s_i, s_{-i}^*)$, which is a contradiction. $\square$

Let S be the resultant strategy space obtained from IESDS and N be the set of all Nash equilibria of a game, then the above proposition essentially means that $N \subseteq S$.

Note that this does not mean that Proposition 1.1.8 establishes the **existence** of Nash equilibria for any game. However, we can indeed conclude a little bit about the uniqueness of Nash equilibrium under some circumstances.

Proposition 1.1.9 ► Uniqueness of Nash Equilibrium from IESDS

Let $G = (S_1, S_2, \dots, S_n, u_1, u_2, \dots, u_n)$ be an n -player normal-form game with finite strategy spaces. If $(s_1^*, s_2^*, \dots, s_n^*) \in S$ are the only strategies which are not eliminated by IESDS, then it is the unique Nash equilibrium.

Proof. The uniqueness is obvious from Proposition 1.1.8, so it suffices to prove that the strategies $(s_1^*, s_2^*, \dots, s_n^*)$ are indeed a Nash equilibrium. Suppose on contrary that it is not a Nash equilibrium, then there exists some s_i^* such that $s_i^* \notin R_i(s_{-i}^*)$. Take some $b_i \in R_i(s_{-i}^*)$, then

$$u_i(b_i, s_{-i}^*) > u_i(s_i^*, s_{-i}^*).$$

However, since b_i is eliminated, it must be strictly dominated by some $t_i \in S_i$ such that $t_i \neq b_i$. Therefore, we have

$$u_i(t_i, s_{-i}^*) > u_i(b_i, s_{-i}^*)$$

because s_{-i}^* are strategies which are not eliminated, but this is a contradiction since t_i is a better response than b_i . □

So far, in our analysis we have been assuming that the players in a game make decisions **deterministically**, i.e., the player always choose a single strategy as the response. However, in real-life scenarios, there could be some random factors involved in decision-making.

Definition 1.1.10 ► Mixed Strategy

Let $G = (S_1, S_2, \dots, S_n, u_1, u_2, \dots, u_n)$ be an n -player normal-form game with finite strategy spaces. Suppose $S_i = \{s_{i1}, s_{i2}, \dots, s_{ik}\}$, then each $s_{ij} \in S_i$ is a **pure strategy** for player i and a probability distribution $p_i := (p_{i1}, p_{i2}, \dots, p_{ik})$ is said to be a **mixed strategy** such that the probability of player i playing s_{ij} is p_{ij} .

As an example, we consider two players, i and j , with M and N pure strategies respectively. Let p_i and p_j be mixed strategies for i and j respectively, then the **expected payoff**

of player i playing any pure strategy s_{ik} is

$$v_i(s_{ik}, p_j) := \sum_{n=1}^N p_{jn} u_i(s_{ik}, s_{jn}).$$

With this, we can compute the **expected payoff** for the mixed strategy p_i as

$$\begin{aligned} v_i(p_i, p_j) &:= \sum_{m=1}^M p_{im} v_i(s_{im}, p_j) \\ &= \sum_{m=1}^M \sum_{n=1}^N p_{im} p_{jn} u_i(s_{im}, s_{jn}). \end{aligned}$$

Similarly, the expected payoff for p_j is

$$v_j(p_j, p_i) = \sum_{n=1}^N \sum_{m=1}^M p_{jn} p_{im} u_j(s_{jn}, s_{im}).$$

Now we can extend the notion of Nash equilibrium to mixed strategies using expected payoffs.

Definition 1.1.11 ► Nash Equilibrium for Mixed Strategies

Let $G := \{S_1, S_2, \dots, S_n, u_1, u_2, \dots, u_n\}$ be a normal-form n -player game, then the probability distributions $(p_1^*, p_2^*, \dots, p_n^*)$ are a **Nash equilibrium** if

$$v_i(p_i^*, \mathbf{p}_{-i}^*) \geq v_i(p_i, \mathbf{p}_{-i}^*)$$

for any probability distribution p_i on S_i .

Intuitively, if a pure strategy s_{ik} is eliminated by IESDS, then any mixed strategy that is good enough probably should not favour playing the strategy. Indeed, we have the following result:

Proposition 1.1.12 ► Zero Probability of Strategies Eliminated by IESDS

Let $(p_1^, p_2^*, \dots, p_n^*)$ be any mixed strategy Nash equilibrium. If s_{ik} is eliminated by IESDS, then $p_{ik}^* = 0$.*

Lastly, combining the two scenarios of Nash equilibria, the following theorem states formally the existence of Nash equilibria:

Theorem 1.1.13 ▶ Nash's Theorem

Let $G := \{S_1, S_2, \dots, S_n, u_1, u_2, \dots, u_n\}$ be a normal-form n -player game with finite strategy spaces, then there exists at least one Nash equilibrium if mixed strategies are allowed.

1.2 Dynamic Games of Complete Information

In static games, the players move independently and the strategy of one player will not affect the context in which another player would make decisions. However, there are many games where this is not true. For example, in a turn-based game, the sequence in which the players move makes a significant impact on the strategy of each player. Furthermore, such games consist of multiple stages and so a player has to observe the past actions in order to predict the best decision.

Note that now, the extent to which a player can observe the past is important as this impacts how much information the player can acquire to make a decision in his or her own profit.

1.2.1 Perfect Information

We will first discuss the first case where the information about the past moves are fully available.

Definition 1.2.1 ▶ Dynamic Games of Complete and Perfect Information

A **dynamic game** of **complete** and **perfect** information is a game where

- the players move in sequence,
- all previous moves are observed before the next move is taken, and
- payoffs are common knowledge to all players.

Such games can be presented with a *game tree* consisting of game states after each possible sequence of moves. A game state can be modelled by an ordered n -tuple $(x_{k1}, x_{k2}, \dots, x_{kn})$ denoting the payoffs off all players after k moves. In this respect, the initial state can be represented by $\mathbf{0}$. Let $G := \{S_1, S_2, \dots, S_n, u_1, u_2, \dots, u_n\}$ be such a dynamic game.

Let $\mathbf{0}$ be the root of the game tree. For each payoff vector $\mathbf{u}_{h-1}$ in height $h - 1$, the h -th player takes a move in the *feasible set* A_h to modify the payoffs to a new payoff vector $\mathbf{u}_h$ and create an edge $s_h := \mathbf{u}_{h-1} \cdot \mathbf{u}_h$. Consider the leaves of the tree formed this way at height n . Note that every strategy is uniquely represented as the $\mathbf{0} \cdot \mathbf{u}$ path for some leaf $\mathbf{u}$, where the projection $\pi_i(\mathbf{u}) = u_i(s_1, s_2, \dots, s_n)$ is the payoff corresponding to the strategy.

The optimal strategy of a dynamic game of complete and perfect information can be obtained from a recursive algorithm.

Technique 1.2.2 ► Backward Induction